

System Design: Real-time Olympic Stats Platform

David Bauch

2024-12-31

Table of contents

System Design	1
Learning Goals	1
Situation	2
Deliverables	2
How the System Works	3
Potential Bottlenecks	5
Scaling Solutions	6
Security Boundaries	7
Security Measures	8
Assumptions	10
Volumes	10
Calculations (Usage-Based)	10
Usage-Based Cost Table (17 days)	11
Enterprise Plan Consideration	11
Summary	12

System Design

Learning Goals

- Design a system that handles high traffic with real-time updates.
- Identify bottlenecks, scaling limits, and security boundaries.
- Estimate costs for peak and moderate traffic.

Situation

An investor asks you to take \$1,000,000 (one million dollars) and deploy a website that hosts realtime stats for all events in the next Winter Olympics including leaderboards by different categories.

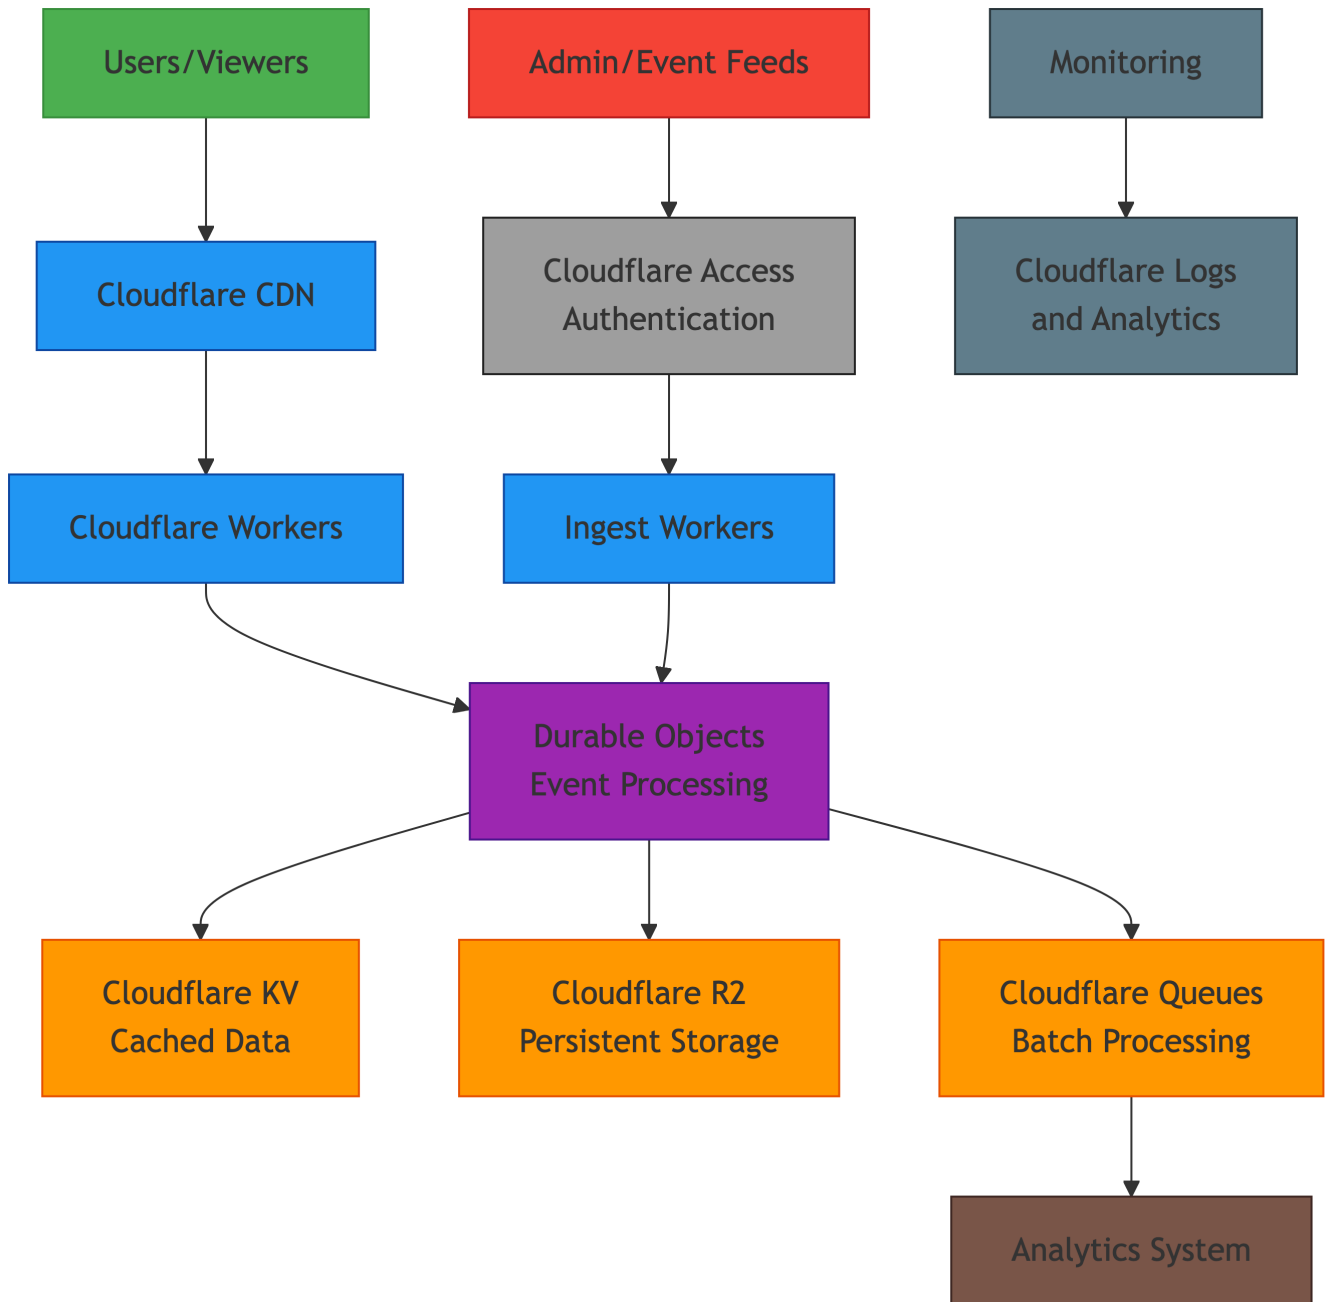
Some quick research tells you you can expect about **360,000,000 requests per day** during **17 days** of the Olympic games.

Deliverables

Use Cloudflare for everything.

1. System Design

- Draw a diagram showing all components and how requests flow.
- Explain how your system handles traffic and scales.



How the System Works

1. User Request Flow

- Users access the Olympic stats website through **Cloudflare CDN**, which serves cached content from the edge locations closest to them
- For dynamic content, requests are handled by **Cloudflare Workers** that run serverless functions at the edge
- Workers fetch data from **Cloudflare KV** (key-value store) for fast, globally-distributed reads

2. Real-time Data Processing

- **Durable Objects** act as the core processing units, with each Olympic event having its own dedicated object
- These objects maintain state and ensure consistency for real-time updates
- When new data arrives (scores, times, etc.), Durable Objects:
 - Process and validate the update
 - Update the leaderboard calculations
 - Store the latest state in **Cloudflare KV** for fast reads
 - Archive data to **Cloudflare R2** (S3-compatible object storage) for persistence

3. Data Ingestion

- **Admin/Event feeds** (trusted sources) send updates through a secure channel
- **Cloudflare Access** provides authentication for these administrative endpoints
- **Ingest Workers** validate and sanitize incoming data before passing it to Durable Objects

4. Scalability Features

- **Automatic scaling:** Cloudflare automatically scales Workers and Durable Objects based on demand
- **Global distribution:** All components run at Cloudflare's edge locations worldwide
- **Batch processing:** Non-critical tasks are queued using **Cloudflare Queues** for asynchronous processing
- **Caching layers:** Multiple caching layers reduce load on processing components

5. Monitoring and Analytics

- **Cloudflare Logs and Analytics** provide real-time monitoring of system performance
- Metrics are collected on request rates, cache hit ratios, and processing latencies
- An integrated **Analytics System** processes batched data for reporting and insights

This architecture handles the high traffic requirements (360M requests/day) by pushing computation to the edge and heavily leveraging caching. The system is designed to scale automatically with demand while maintaining low-latency responses for users worldwide.

2. Bottlenecks & Scaling

- Identify potential bottlenecks (e.g., real-time processing, database writes).
- Explain how your design addresses scaling constraints.

Potential Bottlenecks

1. Durable Objects Serialization

- **Bottleneck:** Each Durable Object processes updates sequentially, which could create a queue if updates arrive faster than they can be processed
- **Mitigation:**
 - Limit the number of concurrent events per Object
 - Use multiple Durable Objects for high-traffic events
 - Implement efficient batching of updates within the Object

2. KV Write Capacity

- **Bottleneck:** Cloudflare KV has rate limits on writes (approximately 1 write per second per key in the free tier, higher in paid tiers)
- **Mitigation:**
 - Batch updates to reduce write frequency
 - Use different keys for different granularities of data
 - Implement exponential backoff for rate limit handling

3. Network Bandwidth for Real-time Updates

- **Bottleneck:** Broadcasting real-time updates to thousands of connected clients simultaneously
- **Mitigation:**
 - Use Cloudflare's global network to distribute the load
 - Implement client-side throttling to limit update frequency
 - Compress data payloads

4. Event Feed Ingestion

- **Bottleneck:** Handling spikes in incoming data from event feeds
- **Mitigation:**
 - Queue incoming updates using Cloudflare Queues
 - Implement rate limiting at the ingest worker level
 - Use auto-scaling ingest workers

Scaling Solutions

1. Horizontal Scaling

- **Approach:** The system scales horizontally by adding more instances of components
- **Implementation:**
 - Cloudflare automatically adds more Worker instances based on demand
 - New Durable Objects are created for new events
 - Additional KV namespaces can be used for very large datasets

2. Geographic Distribution

- **Approach:** All components run at Cloudflare's edge locations worldwide
- **Benefits:**
 - Reduced latency for users
 - Load distribution across multiple data centers
 - Improved fault tolerance

3. Caching Strategy

- **Approach:** Multi-layer caching reduces load on processing components
- **Implementation:**
 - CDN caching for static and semi-static content
 - Worker response caching
 - Client-side caching with appropriate TTLs

4. Asynchronous Processing

- **Approach:** Non-critical operations are processed asynchronously
- **Implementation:**
 - Analytics and reporting via Cloudflare Queues
 - Data archiving to R2 done asynchronously
 - Batch updates to reduce immediate processing load

5. Resource Allocation

- **Approach:** Cloudflare's platform automatically allocates resources based on demand
- **Benefits:**
 - No manual intervention needed for scaling
 - Cost efficiency during low-traffic periods
 - Automatic handling of traffic spikes

3. Security

- Identify security boundaries (e.g., admin endpoints, APIs).
- Explain how security is maintained (encryption, authentication, firewalls).

Security Boundaries

1. Admin Endpoints

- **Boundary:** Ingest APIs for receiving real-time event data
- **Protection:**
 - Isolated in private Worker routes not accessible from the public internet
 - Protected by Cloudflare Access for zero-trust authentication
 - Rate limited to prevent abuse
 - IP allowlisting for trusted event feed sources

2. Public APIs

- **Boundary:** Read-only endpoints serving stats and leaderboard data
- **Protection:**
 - CORS restricted to official domains
 - Rate limited per IP to prevent scraping
 - Cached responses to reduce direct origin hits
 - Content Security Policy (CSP) headers

3. Internal Communication

- **Boundary:** Communication between Workers, Durable Objects, KV, and Queues
- **Protection:**
 - All communication happens within Cloudflare’s internal network
 - Service tokens for inter-service authentication
 - Encrypted at rest and in transit

4. Data Storage

- **Boundary:** Cloudflare KV and R2 storage systems
- **Protection:**
 - Data encrypted at rest using Cloudflare’s encryption
 - Access control through Worker bindings
 - Regular key rotation

Security Measures

1. Encryption

- **In Transit:**
 - All communication uses HTTPS/TLS 1.3
 - API tokens and sensitive data encrypted in transit
- **At Rest:**
 - Cloudflare KV and R2 automatically encrypt data at rest
 - Secrets stored in encrypted environment variables

2. Authentication

- **User Authentication:**
 - Cloudflare Access for admin interfaces
 - Signed URLs with time-limited tokens for premium content
- **Service Authentication:**
 - Service tokens for inter-Worker communication
 - Mutual TLS (mTLS) for critical service-to-service communication
 - HMAC signatures for event feed validation

3. Firewalls and DDoS Protection

- **Cloudflare WAF:**
 - Pre-configured OWASP rules
 - Custom rules for Olympic-specific attack patterns
 - Real-time rule updates
- **DDoS Protection:**
 - Automatic DDoS mitigation at Cloudflare's edge
 - Rate limiting rules for API endpoints
 - Geo-blocking if necessary

4. Monitoring and Auditing

- **Logging:**
 - All access logs sent to Cloudflare Logs
 - Security events flagged and alerted
 - Audit trails for admin actions
- **Monitoring:**
 - Real-time anomaly detection
 - Automated alerts for suspicious activity
 - Regular security scanning

5. Incident Response

- **Preparation:**
 - Pre-defined incident response procedures
 - Backup and rollback mechanisms
 - Regular security assessments
- **Response:**
 - Automated blocking of malicious IPs
 - Emergency key rotation procedures
 - Communication plan for stakeholders

4. Cost Estimation

- Calculate costs for peak load (17 days) and half load (17 days).
- Include a table of technologies you'll use and associated costs, showing your calculations.

Assumptions

- **Traffic:** Peak 360M requests/day; Half 180M requests/day; both for 17 days.
- **Cache hit:** 99.9% of reads served from CDN cache without Worker execution.
- **Worker/Durable Object on cache miss:** 0.1% of reads invoke Workers and 1 Durable Object op.
- **Ingest writes:** Peak period 20 updates/sec for 12h/day; Half period 10 updates/sec for 12h/day.
- **Payload sizes:** Read response 20 KB avg; write event 2 KB archived to R2.
- **KV writes:** 0.5 KV write per ingest update on average (coalesced snapshots).
- **Queues:** 1 enqueue + 1 dequeue per ingest update.
- **Pricing (public list, rounded):** Workers \$0.30/million req; Durable Objects \$0.15/million req; KV reads \$0.50/million; KV writes \$5.00/million; Queues \$0.40/million ops; R2 storage \$0.015/GB-month. CDN egress is typically covered under Enterprise; usage-priced egress not included here.

Volumes

- **Reads (Peak):** $360,000,000/\text{day} \times 17 = 6.12\text{B}$ total reads; 0.1% misses $\rightarrow 6.12\text{M}$ Worker+Durable Object ops.
- **Reads (Half):** $180,000,000/\text{day} \times 17 = 3.06\text{B}$ total reads; 0.1% misses $\rightarrow 3.06\text{M}$ Worker+Durable Object ops.
- **Writes (Peak):** $20/\text{sec} \times 43,200 \text{ sec/day} \times 17 = 14.688\text{M}$ updates.
- **Writes (Half):** $10/\text{sec} \times 43,200 \text{ sec/day} \times 17 = 7.344\text{M}$ updates.
- **KV Reads:** equal to read misses.
- **KV Writes (Peak/Half):** $0.5 \times \text{updates} \rightarrow 7.344\text{M} / 3.672\text{M}$.
- **Queues Ops (Peak/Half):** $2 \times \text{updates} \rightarrow 29.376\text{M} / 14.688\text{M}$.
- **R2 Data Added (Peak/Half):** $\text{updates} \times 2 \text{ KB} \rightarrow \sim 28.7 \text{ GB} / \sim 14.3 \text{ GB}$.

Calculations (Usage-Based)

- **Workers requests:** (misses + writes)
 - Peak: $6.12\text{M} + 14.688\text{M} = 20.808\text{M} \rightarrow 20.808 \times \$0.30 = \$6.24$
 - Half: $3.06\text{M} + 7.344\text{M} = 10.404\text{M} \rightarrow 10.404 \times \$0.30 = \$3.12$
- **Durable Objects requests:** \sim writes count
 - Peak: $14.688\text{M} \rightarrow 14.688 \times \$0.15 = \$2.20$
 - Half: $7.344\text{M} \rightarrow 7.344 \times \$0.15 = \$1.10$
- **KV reads:**
 - Peak: $6.12\text{M} \rightarrow 6.12 \times \$0.50 = \$3.06$

– Half: $3.06\text{M} \rightarrow 3.06 \times \$0.50 = \$1.53$

- **KV writes:**

– Peak: $7.344\text{M} \rightarrow 7.344 \times \$5.00 = \$36.72$

– Half: $3.672\text{M} \rightarrow 3.672 \times \$5.00 = \$18.36$

- **Queues ops:**

– Peak: $29.376\text{M} \rightarrow 29.376 \times \$0.40 = \$11.75$

– Half: $14.688\text{M} \rightarrow 14.688 \times \$0.40 = \$5.88$

- **R2 storage (1 month):**

– Peak: $\sim 30\text{ GB} \times \$0.015 = \$0.45$

– Half: $\sim 15\text{ GB} \times \$0.015 = \$0.23$

Usage-Based Cost Table (17 days)

Service	Unit Price	Peak Quantity	Peak Cost	Half Quantity	Half Cost
Workers requests	\$0.30 / 1M	20.808M	\$6.24	10.404M	\$3.12
Durable Objects req	\$0.15 / 1M	14.688M	\$2.20	7.344M	\$1.10
KV reads	\$0.50 / 1M	6.12M	\$3.06	3.06M	\$1.53
KV writes	\$5.00 / 1M	7.344M	\$36.72	3.672M	\$18.36
Queues ops	\$0.40 / 1M	29.376M	\$11.75	14.688M	\$5.88
R2 storage	\$0.015 / GB-mo	30 GB	\$0.45	15 GB	\$0.23
CDN egress	Included in Enterprise	122 TB	—	61 TB	—
TOTAL (usage fees only)			\$60.42		\$30.22

Enterprise Plan Consideration

- Production traffic at Olympic scale typically requires a Cloudflare Enterprise contract (WAF, Bot Management, SLA, Logpush, priority support, custom limits).

- For budgeting, assume an Enterprise engagement for the Olympic window at a flat fee. Example placeholder: **\$200,000** for 1–2 months. This is an assumption and must be vendor-quoted.

Summary

- **Usage-based fees are minimal** relative to traffic due to high cache hit rates and per-million pricing (tens of dollars for the compute/storage pieces).
- **Primary cost driver is the Enterprise contract**, which provides the necessary protections and support at this scale.
- With a \$1,000,000 budget, this design leaves ample headroom for operations, monitoring, contingency, and team costs during the event window.